

R Notebook

This is an R Markdown Notebook. When you execute code within the notebook, the results appear beneath the code.

Try executing this chunk by clicking the *Run* button within the chunk or by placing your cursor inside it and pressing *Ctrl+Shift+Enter*.

```
health <- read.csv("C:\\Users\\srava\\Documents\\Ryerson university\\Capstone project\\Maternal Health I
health2 <- health;
health2[health2$RiskLevel == "high risk", "RiskLevel"] <- 2
health2[health2$RiskLevel == "low risk", "RiskLevel"] <- 0
health2[health2$RiskLevel == "mid risk", "RiskLevel"] <- 1
head(health2)
```

```
##   i..Age SystolicBP DiastolicBP   BS BodyTemp HeartRate RiskLevel
## 1    25        130          80 15.00        98        86         2
## 2    35        140          90 13.00        98        70         2
## 3    29         90          70  8.00       100        80         2
## 4    30        140          85  7.00        98        70         2
## 5    35        120          60  6.10        98        76         0
## 6    23        140          80  7.01        98        70         2
```

```
str(health2)
```

```
## 'data.frame':   1014 obs. of  7 variables:
##  $ i..Age      : int  25 35 29 30 35 23 23 35 32 42 ...
##  $ SystolicBP  : int  130 140 90 140 120 140 130 85 120 130 ...
##  $ DiastolicBP : int   80 90 70 85 60 80 70 60 90 80 ...
##  $ BS          : num  15 13 8 7 6.1 7.01 7.01 11 6.9 18 ...
##  $ BodyTemp    : num  98 98 100 98 98 98 98 102 98 98 ...
##  $ HeartRate   : int   86 70 80 70 76 70 78 86 70 70 ...
##  $ RiskLevel   : chr   "2" "2" "2" "2" ...
```

Add a new chunk by clicking the *Insert Chunk* button on the toolbar or by pressing *Ctrl+Alt+I*.

```
health2$RiskLevel<- as.numeric(as.character(health2$RiskLevel))
## Shuffling the data to make sure all the classes are included without bias
shuffle_index <- sample(1:nrow(health2))
head(shuffle_index)
```

```
## [1] 433 118  58 877 395 664
```

```
health2 <- health2[shuffle_index, ]
```

```
head(health2)
```

```
##      i..Age SystolicBP DiastolicBP   BS BodyTemp HeartRate RiskLevel
## 433    40      140      100 13.0      101      66      2
## 118    55      140      100 18.0      98      90      2
## 58     42      120      80  6.4      98      70      0
## 877    27      120      70  6.8      98      77      0
## 395    19      120      80  7.0      98      70      1
## 664    15       90      60  6.0      98      80      0
```

```
## Splitting the dataset into training and test in 1:4 ratio
```

```
create_train_test <- function(data, size = 0.8, train = TRUE) {
  n_row = nrow(data)
  total_row = size * n_row
  train_sample <- 1: total_row
  if (train == TRUE) {
    return (data[train_sample, ])
  } else {
    return (data[-train_sample, ])
  }
}
```

```
data_train <- create_train_test(health2, 0.8, train = TRUE)
data_test <- create_train_test(health2, 0.8, train = FALSE)
```

```
## Dimensions of test and training data
dim(data_train)
```

```
## [1] 811  7
```

```
dim(data_test)
```

```
## [1] 203  7
```

```
## Decision tree model
```

```
install.packages("rpart.plot", repos = "http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/srava/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
```

```
## package 'rpart.plot' successfully unpacked and MD5 sums checked
```

```
##
```

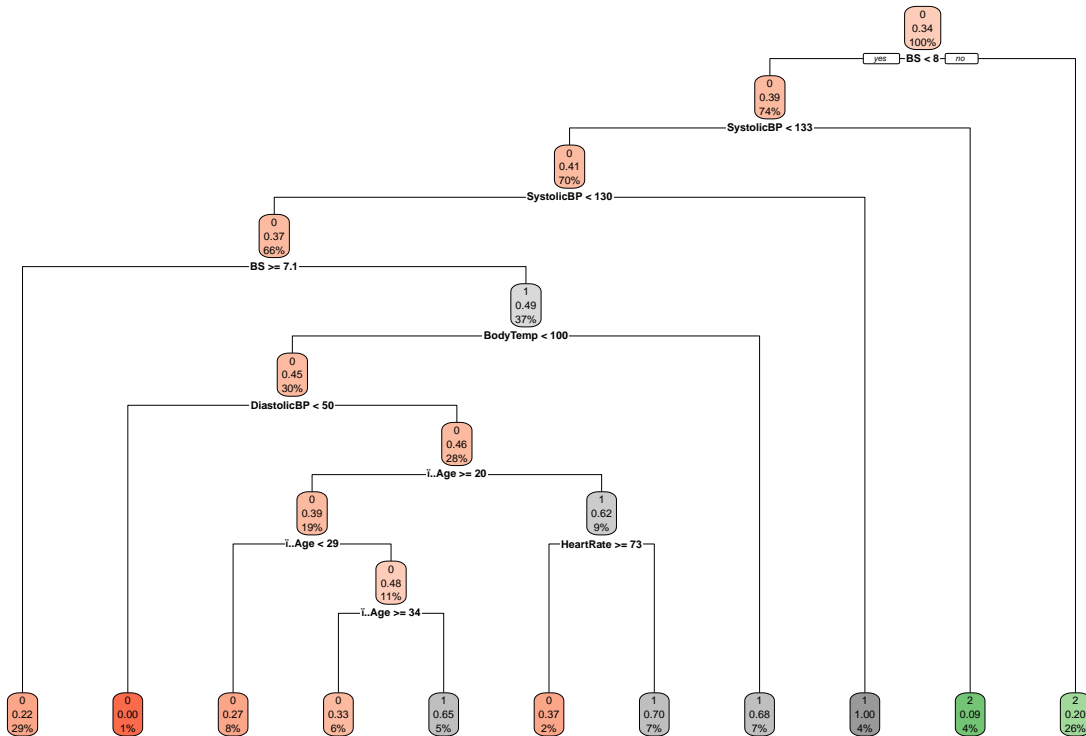
```
## The downloaded binary packages are in
```

```
## C:\Users\srava\AppData\Local\Temp\RtmpW8mwQi\downloaded_packages
```

```
library(rpart)
library(rpart.plot)
fit <- rpart(RiskLevel~., data = data_train, method = 'class')
rpart.plot(fit, extra = 106)
```

```
## Warning: extra=106 but the response has 3 levels (only the 2nd level is
## displayed)
```

■ 0
 ■ 1
 ■ 2



```
predict_unseen <- predict(fit, data_test, type = 'class')
```

```
table_con <- table(data_test$RiskLevel, predict_unseen)
table_con
```

```
##      predict_unseen
##      0  1  2
##  0 75 14  1
##  1 21 23 16
##  2  6  1 46
```

```
accuracy_Test <- sum(diag(table_con)) / sum(table_con)
```

```
print(paste('Accuracy for test', accuracy_Test))
```

```
## [1] "Accuracy for test 0.70935960591133"
```

```
# Installing Packages
install.packages("e1071", repos = "http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/srava/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
```

```
## package 'e1071' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\srava\AppData\Local\Temp\RtmpW8mwQi\downloaded_packages
```

```
install.packages("caTools",repos ="http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/srava/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
```

```
## package 'caTools' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\srava\AppData\Local\Temp\RtmpW8mwQi\downloaded_packages
```

```
install.packages("class",repos ="http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/srava/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
```

```
## package 'class' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\srava\AppData\Local\Temp\RtmpW8mwQi\downloaded_packages
```

```
# Loading package
```

```
library(e1071)
library(caTools)
library(class)
```

```
#Naive Bayes
```

```
# Installing Packages
```

```
install.packages("caret",repos ="http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/srava/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
```

```
## package 'caret' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\srava\AppData\Local\Temp\RtmpW8mwQi\downloaded_packages
```

```
# Loading package
```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```

# Splitting data into train
# and test data
split <- sample.split(health2, SplitRatio = 0.7)
train_cl <- subset(health2, split == "TRUE")
test_cl <- subset(health2, split == "FALSE")

# Feature Scaling
train_scale <- scale(train_cl[, 1:6])
test_scale <- scale(test_cl[, 1:6])

# Fitting Naive Bayes Model
# to training dataset
set.seed(120) # Setting Seed
classifier_cl <- naiveBayes(RiskLevel ~ ., data = train_cl)
classifier_cl

```

```

##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##      0      1      2
## 0.4034483 0.3465517 0.2500000
##
## Conditional probabilities:
##   i..Age
## Y      [,1]      [,2]
## 0 26.70940 12.79489
## 1 28.02985 12.31987
## 2 34.95172 12.54178
##
##   SystolicBP
## Y      [,1]      [,2]
## 0 106.6581 15.41203
## 1 112.8358 15.26984
## 2 123.2966 20.90293
##
##   DiastolicBP
## Y      [,1]      [,2]
## 0 72.62821 12.98012
## 1 73.92537 11.12427
## 2 84.57241 14.38737
##
##   BS
## Y      [,1]      [,2]
## 0  7.174487 0.5761944
## 1  7.747861 2.1370678
## 2 11.998828 4.1668667
##
##   BodyTemp

```

```
## Y      [,1]      [,2]
## 0 98.37778 1.128569
## 1 98.83085 1.459878
## 2 99.01793 1.652257
##
## HeartRate
## Y      [,1]      [,2]
## 0 72.74786 8.066282
## 1 74.10448 6.667386
## 2 76.07586 8.434455
```

```
# Predicting on test data
y_pred <- predict(classifier_cl, newdata = test_cl)
# Confusion Matrix
cm <- table(test_cl$RiskLevel, y_pred)
cm
```

```
##      y_pred
##      0    1    2
## 0 154   14    4
## 1  88   31   16
## 2  25   19   83
```

```
confusionMatrix(cm)
```

```
## Confusion Matrix and Statistics
##
##      y_pred
##      0    1    2
## 0 154   14    4
## 1  88   31   16
## 2  25   19   83
##
## Overall Statistics
##
##              Accuracy : 0.6175
##              95% CI : (0.57, 0.6634)
##      No Information Rate : 0.6152
##      P-Value [Acc > NIR] : 0.4818
##
##              Kappa : 0.4032
##
## Mcnemar's Test P-Value : 6.489e-15
##
## Statistics by Class:
##
##              Class: 0 Class: 1 Class: 2
## Sensitivity          0.5768 0.48438 0.8058
## Specificity          0.8922 0.71892 0.8671
## Pos Pred Value       0.8953 0.22963 0.6535
## Neg Pred Value       0.5687 0.88963 0.9349
## Prevalence           0.6152 0.14747 0.2373
## Detection Rate       0.3548 0.07143 0.1912
```

```
## Detection Prevalence    0.3963  0.31106  0.2926
## Balanced Accuracy      0.7345  0.60165  0.8364
```

Model Evaluation

```
confusionMatrix(cm)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      y_pred
```

```
##      0    1    2
```

```
## 0 154   14    4
```

```
## 1  88   31   16
```

```
## 2  25   19   83
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##              Accuracy : 0.6175
```

```
##              95% CI : (0.57, 0.6634)
```

```
##      No Information Rate : 0.6152
```

```
##      P-Value [Acc > NIR] : 0.4818
```

```
##
```

```
##              Kappa : 0.4032
```

```
##
```

```
## McNemar's Test P-Value : 6.489e-15
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##              Class: 0 Class: 1 Class: 2
```

```
## Sensitivity      0.5768  0.48438  0.8058
```

```
## Specificity      0.8922  0.71892  0.8671
```

```
## Pos Pred Value   0.8953  0.22963  0.6535
```

```
## Neg Pred Value   0.5687  0.88963  0.9349
```

```
## Prevalence       0.6152  0.14747  0.2373
```

```
## Detection Rate   0.3548  0.07143  0.1912
```

```
## Detection Prevalence 0.3963  0.31106  0.2926
```

```
## Balanced Accuracy 0.7345  0.60165  0.8364
```

knn algorithm

```
install.packages("e1071",repos ="http://cran.us.r-project.org")
```

```
## Warning: package 'e1071' is in use and will not be installed
```

```
install.packages("caTools",repos ="http://cran.us.r-project.org")
```

```
## Warning: package 'caTools' is in use and will not be installed
```

```
install.packages("class",repos ="http://cran.us.r-project.org")
```

```
## Warning: package 'class' is in use and will not be installed
```

```

library(caret)
library(e1071)
library(caTools)
library(class)
# Splitting the sample
split <- sample.split(health2, SplitRatio = 0.7)
train_cl <- subset(health2, split == "TRUE")
test_cl <- subset(health2, split == "FALSE")

# Splitting the sample
split <- sample.split(health2, SplitRatio = 0.7)
train_cl <- subset(health2, split == "TRUE")
test_cl <- subset(health2, split == "FALSE")

# Feature Scaling
train_scale <- scale(train_cl[, 1:6])
test_scale <- scale(test_cl[, 1:6])

classifier_knn <- knn(train = train_scale, test = test_scale,
                      cl = train_cl$RiskLevel,
                      k = 1)
classifier_knn

##      [1] 2 0 1 1 1 2 0 0 1 0 1 1 2 2 2 2 0 2 2 1 1 0 2 0 1 1 1 0 2 0 2 1 1 1 0 1 2
##     [38] 1 1 1 0 0 0 2 1 1 2 2 1 1 1 1 1 0 2 1 0 0 2 1 1 2 2 0 2 0 1 0 0 1 1 0 0
##     [75] 0 1 1 0 0 2 0 1 1 1 0 0 2 1 1 2 1 0 0 1 1 1 2 1 1 0 1 1 2 0 1 0 2 0 1 0 1
##    [112] 2 0 1 0 1 2 1 1 0 2 1 2 2 1 2 2 1 2 1 2 0 1 2 2 1 2 1 1 2 2 1 1 1 2 0 1 0
##    [149] 0 0 1 0 2 1 1 0 0 1 2 1 1 0 0 0 1 1 2 1 1 0 0 1 0 2 0 1 2 1 0 2 1 1 1 0 1
##    [186] 0 1 2 0 0 1 0 1 2 2 1 1 0 0 1 0 1 0 2 0 2 0 1 0 2 1 0 1 1 0 1 1 1 1 0 0 1
##    [223] 2 2 0 2 0 0 0 0 1 0 1 1 1 0 0 2 1 1 1 0 0 0 1 2 1 1 0 2 0 1 2 2 1 1 2 2 0
##    [260] 0 0 0 1 2 2 0 0 1 1 2 1 1 0 0 0 2 1 0 1 1 1 0 1 1 2 0 1 2 2 2 1 1 1 0 0 0
##    [297] 0 1 0 2 1 1 2 1 0 0 1 0 2 1 0 2 0 0 0 0 1 0 0 2 1 1 1 1 2 0 0 0 2 1 1 0 2
##    [334] 0 1 2 2 1 1 1 0 2 1 0 0 1 2 0 2 1 2 0 0 2 0 0 2 0 0 1 2 1 1 2 0 2 1 1 0 0
##    [371] 0 1 2 1 1 0 2 1 2 1 1 0 1 0 2 1 2 2 1 2 1 2 0 2 2 1 0 1 0 0 1 2 0 2 0 0 0
##    [408] 0 1 2 0 1 2 0 0 1 0 0 2 2 2 1 2 1 0 0 0 2 2 0 1 0 0 1 2
## Levels: 0 1 2

cm <- table(test_cl$RiskLevel, classifier_knn)
cm

##      classifier_knn
##           0      1      2
## 0 121    40      9
## 1   25   121      8
## 2    6    11    94

#Model Evaluation - Choosing K
# Calculate out of Sample error
misClassError <- mean(classifier_knn != test_cl$RiskLevel)
print(paste('Accuracy =', 1-misClassError))

## [1] "Accuracy = 0.772413793103448"

```



```
# K = 3
classifier_knn <- knn(train = train_scale,
                      test = test_scale,
                      cl = train_cl$RiskLevel,
                      k = 3)
misClassError <- mean(classifier_knn != test_cl$RiskLevel)
print(paste('Accuracy =', 1-misClassError))
```

```
## [1] "Accuracy = 0.671264367816092"
```

```
# K = 15
classifier_knn <- knn(train = train_scale,
                      test = test_scale,
                      cl = train_cl$RiskLevel,
                      k = 15)
misClassError <- mean(classifier_knn != test_cl$RiskLevel)
print(paste('Accuracy =', 1-misClassError))
```

```
## [1] "Accuracy = 0.67816091954023"
```

```
# K = 19
classifier_knn <- knn(train = train_scale,
                      test = test_scale,
                      cl = train_cl$RiskLevel,
                      k = 1)
misClassError <- mean(classifier_knn != test_cl$RiskLevel)
print(paste('Accuracy =', 1-misClassError))
```

```
## [1] "Accuracy = 0.781609195402299"
```

When you save the notebook, an HTML file containing the code and output will be saved alongside it (click the *Preview* button or press *Ctrl+Shift+K* to preview the HTML file).

The preview shows you a rendered HTML copy of the contents of the editor. Consequently, unlike *Knit*, *Preview* does not run any R code chunks. Instead, the output of the chunk when it was last run in the editor is displayed.